PRE-PRODUCT BUILD ROUND

A marketplace for AI employees you keep.

A planned marketplace where creators sell standardized AI employee packages. Businesses will buy directly or through their AI harness and chosen model.

Shehryar Ali | USD 250,000 target

The tool changes. The job should not.

A repeatable business role should not need rebuilding every time a team switches tools.

WHY NOW

Agent Skills is an open format for reusable expertise across compatible tools.

The opportunity: help buyers choose a tested, maintained role instead of assembling the pieces themselves.

Buy once. Keep the copy. Move with it.

1. Discover

Discover a tested role, yourself or through your AI harness and chosen model.

2. Purchase

Buy a licensed copy. Harness checkout needs permission, a spend cap and a receipt.

3. Keep working

Install in a supported tool. Carry approved context. Buy updates when useful.

Customer control is the product promise. Buying is not execution.
No mandatory marketplace runtime. No subscription to retain the copy. Credentials stay with the buyer.

Prototype benchmarks show technical progress.

WordPress Developer

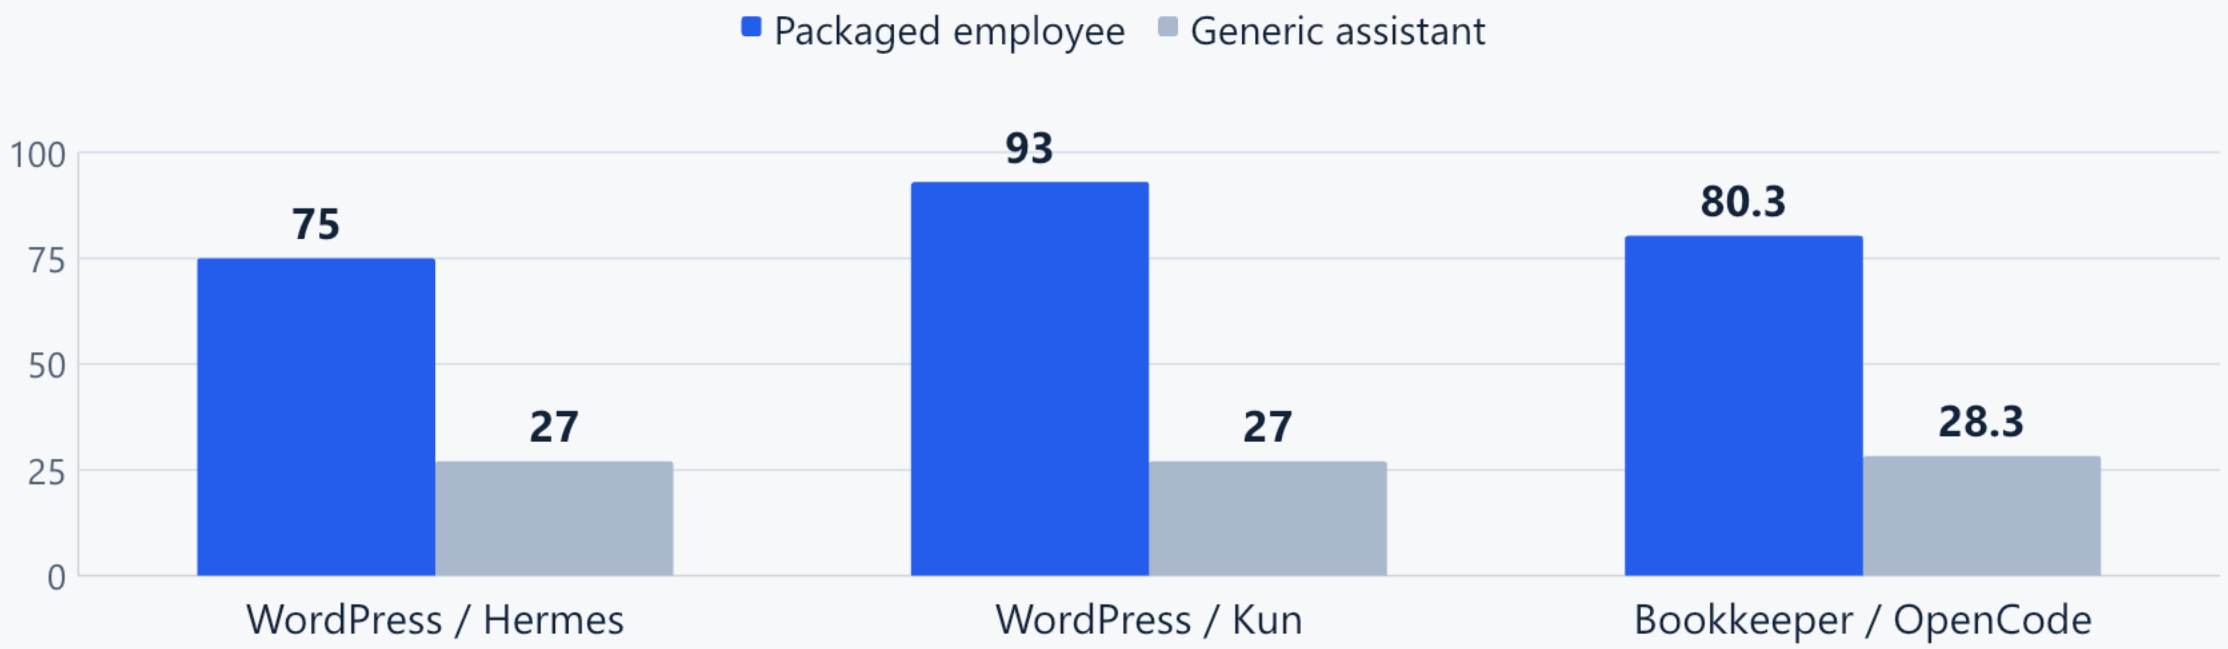
A role prototype with controlled task benchmarks and tested context retrieval across selected environments.

Small Business Bookkeeper

A second role prototype with paired task benchmarks. Further native-environment coverage is still required.

These are technical proofs, not a live marketplace or customer traction.

Higher task scores in controlled benchmarks.



93-point rubric. Three paired runs per series.
Kun generic arm: one timeout.
Controlled tasks, not customer outcomes.

All employee latency targets were missed.
Not customer productivity or universal compatibility.

Portability needs evidence, not promises.

Approved context retrieved

Identity and selected preferences were recovered on tested Hermes, Kun and Claude surfaces.

Credentials were excluded. Personality persistence and every-tool compatibility are not established.

Package scans completed

WordPress: 59 low-severity findings
Bookkeeper: 9 low-severity findings

No medium, high or critical findings in those scans.

Manual review remains required

Bookkeeper scan coverage was incomplete. These are not security certifications.

Start with software operators, then expand.

Recruit supply

Target: 5 independent creator partners.

Offer onboarding, compatibility feedback and a limited founding commission.

Creators bring role expertise and their own audiences.

Activate buyers

Target: 10 buyer design partners from small software and service teams.

Watch real installs and repeat use.

Expand only when buyers get value and creators can maintain quality.

Portability alone is not the moat.

What already exists

Open Agent Skills and community repositories make reusable instructions portable.

A free component is a real alternative. Buyers can assemble their own stack.

What the marketplace must earn

Discovery, purchase rights, tested packages and buyer-controlled context through human UI or AI harness.

Same customers. Two interfaces.

Moat hypothesis: trusted supply and repeat use.

A transaction model aligned with creators.

85% / 15%

Creator / platform share
of net receipts.

Illustration: a USD 99 purchase yields about USD 14.37 to the platform after assumed processing.

Working model, not final terms.

90% / 10%

Proposed founding offer

First 25 qualified creators.
12 months or first USD 10,000 in net receipts, whichever comes first.

Optional maintenance is separate from the purchased copy.

A founder with delivery proof.

Shehryar Ali

Founder and CEO,
S.T.A.R.S Private Limited, Lahore.
Software and game development.

Sole founder of Agent Marketplace.
Full-time from first funding.

Expert-Vetted

Upwork top 1% | 100% Job Success
Personal earnings: $550K Upwork,
$220K Freelancer.

Industry awards

Presented by Pakistan's President:
Best Software House, 2022
Best IT Services, 2023

Six-month target: build funding to buyer pilot.

Build

Human + harness checkout. Permissions, spend caps, receipts and package delivery.

Onboard

5 creator partners.
10 buyer partners.
3 accepted packages.

Validate

5 buyer installs.
3 repeat users.
Test both purchase paths.

Pilot targets, not current traction or a scale forecast.
Both purchase paths are launch scope. Added engineering is not yet re-costed.

USD 250K target. USD 200K core build.

Six-month allocation	USD thousands
Delivery team	86.9
Equipment ceiling	45.0
Engineering software	18.0
Legal and formation	15.0
Infrastructure	15.0
Pilot validation	8.0
Contingency	12.1

Six-month planning case
Founder + architect + QA

USD 250K adds reserve. USD 100K is a solo fallback.

Human + harness checkout costs still need validation.